
Conditional Neural Processes

2020. 02. 13

SeulGi Hong

seulgihong@kaist.ac.kr

Contents

- Introduction
- Concept
- Method
- Experimental Results
- TODO

Combine the benefits of **Stochastic Processes** and **Deep Neural Networks**

- Neural Network
 - structured format
 - trained by gradient descent
- Stochastic Process (e.g. GP)
 - defines the distribution of functions

What is CNP?

CNP is a conditional distribution over functions

which is trained to model the **empirical** conditional distributions of functions $f \sim P$

Concept of Conditional Neural Process Q_θ

Define **conditional distribution over functions** given a set of observations.

- observations are **parameterized** by the **NNs**
- with fixed dimension

$$Q_\theta(f(x_i)|O, x_i) = Q(f(x_i)|\phi_i)$$

How?

⇒ Encoder-Decoder Model

$$\begin{aligned} r_i &= h_\theta(x_i, y_i) & \forall (x_i, y_i) \in O & \quad \text{Encoder} \\ r &= r_1 \oplus r_2 \oplus \dots \oplus r_{n-1} \oplus r_n & & \quad \text{Aggregator} \\ \phi_i &= g_\theta(x_i, r) & \forall (x_i) \in T & \quad \text{Decoder} \end{aligned}$$

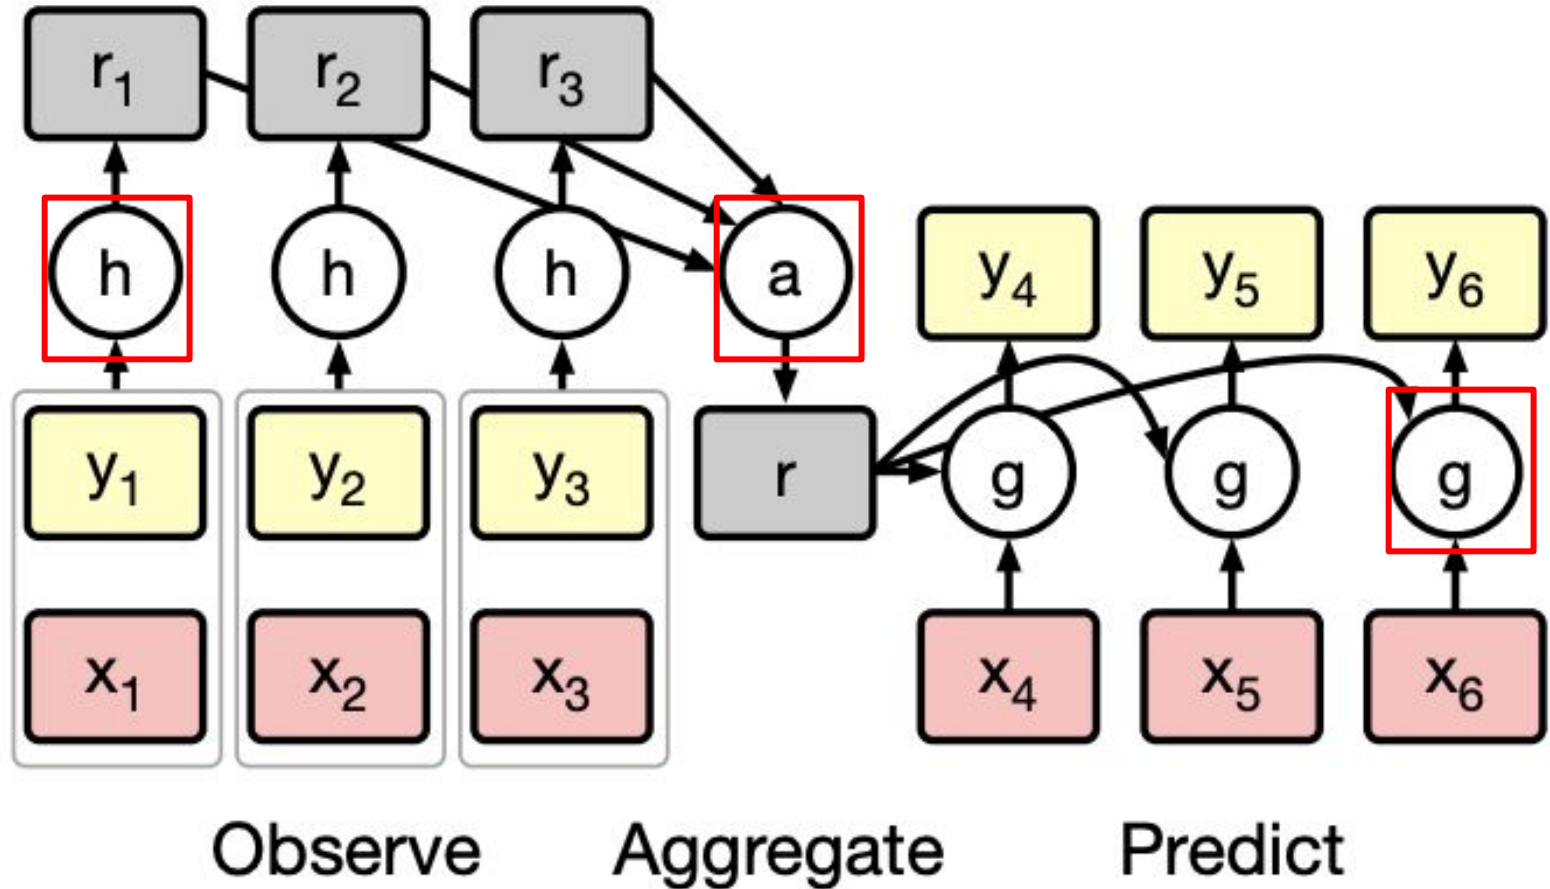
Concept of CNPs

- Shifts the burden of imposing **prior knowledge** from an analytic prior to empirical data
- Invariant under permutations of its inputs.
Specifically, we use factored distribution in this paper

$$Q_{\theta}(f(T)|O, T) = \prod_{x \in T} Q_{\theta}(f(x)|O, x).$$

Method

c Our Model



h function

- input $[N, x_dim]$ and $[N, y_dim]$; **concat**
- output $[N, r_dim]$
- any operation to extract features (e.g. MLP, convolutions...)

$O(N)$

aggregator

- input $[N, r_dim] \rightarrow$ output $[r_dim]$
- commutative operation (e.g. average)
- **compression of observed knowledge**
(just think about mean in GP)

$O(1)$

g function & model output

g function

- input $[r_dim]$ and $[M, x_dim]$
- output $[M, \phi_dim]$; parameterized form

Regression task

- use ϕ to parameterize the **mean** and **variance** of Gaussian distribution

Classification task

- use ϕ to parameterize the **logits of the class probability p_c**

$O(M)$

Loss : Negative conditional log probability

$$L(\theta) = -E_{f \sim P} [E_N [\log Q_{\theta}(\{y_i\}_{i=0}^{n-1} | O_N, \{x_i\}_{i=0}^{n-1})]]$$

Randomly chosen subset of O $O_N = \{(x_i, y_i)\}_{i=0}^N \subset O$

Set of Observations

$$O = \{(x_i, y_i)\}_{i=0}^{n-1}$$

Encoder

- input : context x , y pair (**part of the training data**)
- output : context vector r




Decoder

- input : observation x (**all of the training data**) & context vector r
- output : mean & variance of all observation x

Encoder

- input : context x , y pair (**all of the training data**)
- output : context vector r



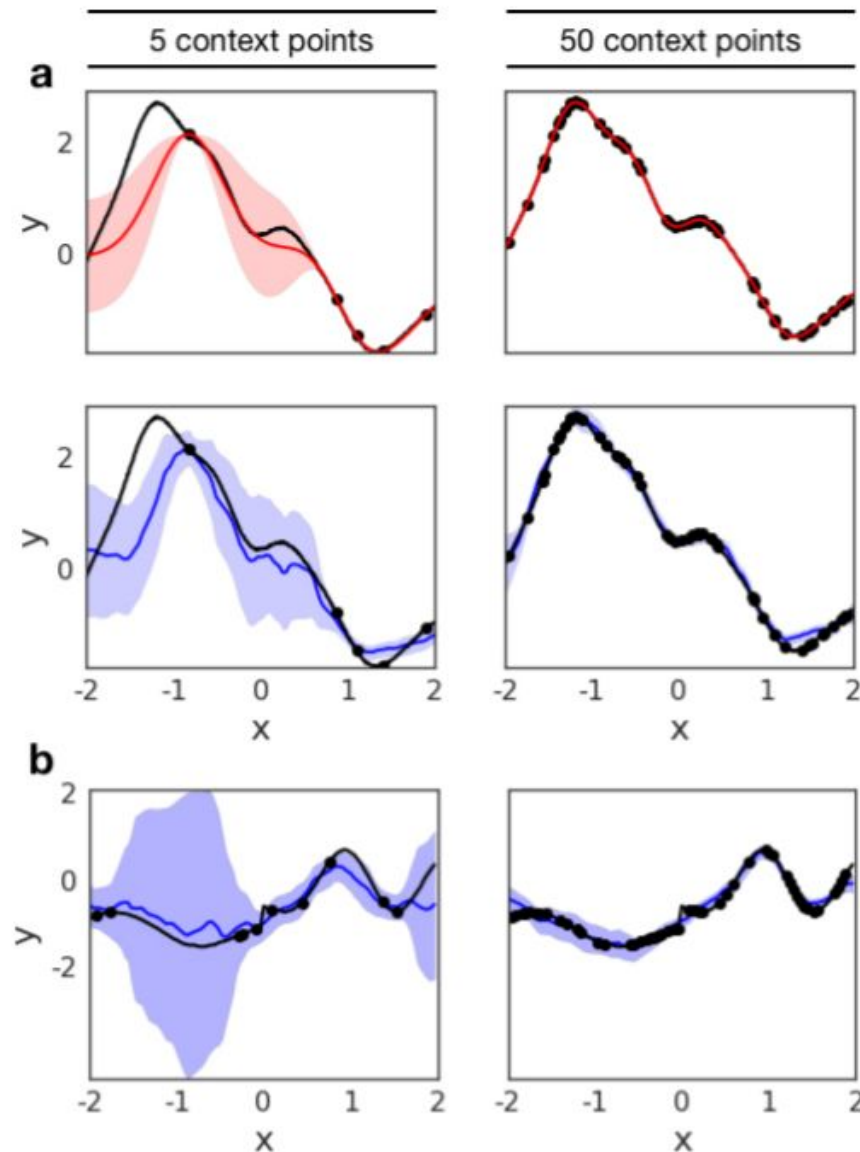
Decoder

- input : observation x (**train + test data**) & context vector r
- output : mean & variance of all observation x

Experimental Results

GP vs CNP

1D Regression

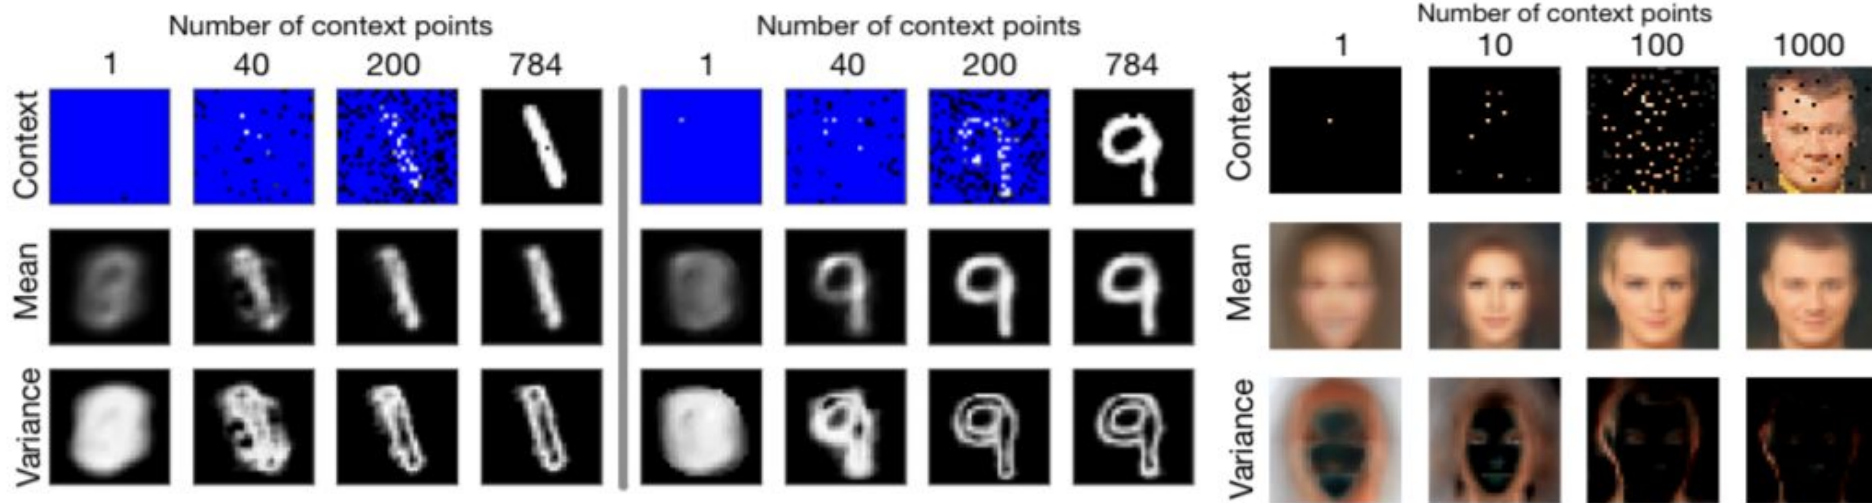


Experimental Results

2D Regression (image completion)

MNIST

CelebA



| # | Random Context | | | Ordered Context | | |
|-----|----------------|--------------|--------------|-----------------|--------------|--------------|
| | 10 | 100 | 1000 | 10 | 100 | 1000 |
| kNN | 0.215 | 0.052 | 0.007 | 0.370 | 0.273 | 0.007 |
| GP | 0.247 | 0.137 | 0.001 | 0.257 | 0.220 | 0.002 |
| CNP | 0.039 | 0.016 | 0.009 | 0.057 | 0.047 | 0.021 |

Experimental Results

Classification on Omniglot

→ achieving comparable performance with inexpensive computational cost



| | 5-way Acc | | 20-way Acc | | Runtime |
|------|--------------|--------------|--------------|--------------|----------------------|
| | 1-shot | 5-shot | 1-shot | 5-shot | |
| MANN | 82.8% | 94.9% | - | - | $\mathcal{O}(nm)$ |
| MN | 98.1% | 98.9% | 93.8% | 98.5% | $\mathcal{O}(nm)$ |
| CNP | 95.3% | 98.5% | 89.9% | 96.8% | $\mathcal{O}(n + m)$ |

Table 2. Classification results on Omniglot. Results on the same task for MANN (Santoro et al., 2016), and matching networks (MN) (Vinyals et al., 2016) and CNP.

TODO

Implement CNP for 1D Regression task

- 1D regression task
- (optional)
2D regression task (image completion)



Link : https://github.com/sghong977/week6_CNP

Conditional Neural Processes implementation

[Manage topics](#)

2 commits

1 branch

Branch: master ▼

New pull request

sghong977 initial commit

TODO

regression1D

regression2D

.gitignore

README.md

TODO

```
# ----- TRAIN -----  
for epoch in range(epochs):  
    nll_loss = 0.0  
    optimizer.zero_grad()
```

```
# TODO -- Train model
```

Train

```
# -----
```

```
if epoch % print_step == 0:  
    print('Epoch', epoch, ': nll loss', nll_loss.item())  
    #dot = make_dot(mean)  
    #dot.render("model.png")  
  
    nll_loss.backward()  
    optimizer.step()
```

```
print("final loss : nll loss", nll_loss.item())  
result_mean, result_var = None, None
```

```
# TODO -- Calculate result_mean and result_var using your model
```

Validation

```
# -----
```

```
draw_graph(x_test, y_test, x_train, y_train, result_mean, np.sqrt(result_var))
```

```
# TODO -- Implement CNP model for 1d regression
```

```
class CNPs(nn.Module):
```

```
    def __init__(self, encoder, decoder):  
        super().__init__()  
        self.encoder = encoder  
        self.decoder = decoder
```

```
    # Output : mean, variance
```

```
    def forward(self, x_ctx, y_ctx, x_obs):  
        pass
```

```
class Encoder(nn.Module):
```

```
    def __init__(self, x_dim=1, y_dim=1, r_dim=128):  
        super().__init__()  
        pass
```

```
    def forward(self, x, y):  
        pass
```

```
class Decoder(nn.Module):
```

```
    def __init__(self, x_dim, r_dim, y_dim):  
        super().__init__()  
        pass
```

```
    def forward(self, x, r):  
        pass
```

Model